

RETAILMART V3 • SQL

Conditional Logic & Derived Columns



WhatsApp Announcement

CONDITIONAL LOGIC & DERIVED COLUMNS

What if your query needs to make DECISIONS? 'If salary > 50000, label as Senior. If NULL, show N/A. Convert text to number.' Today you learn CASE WHEN, COALESCE, NULLIF, and CAST — the decision-making tools of SQL.

Today's Focus:

- CASE WHEN — IF-ELSE logic inside SQL queries
- COALESCE — replace NULLs with defaults
- NULLIF — create NULL conditionally
- CAST & :: — convert between data types

After today, your queries can think and decide!

— Siraj Sir

#Day7 #CASEWHEN #PostgreSQL

TODAY'S AGENDA (2 Hours)

Section	Time	Topics
Quick Recap — earlier sessions	5 mins	String & Date functions from Scalar Functions, Scalar functions = transform one value
Part 1: CASE WHEN — SQL's IF-ELSE	50 mins	Simple CASE (value matching — like switch), Searched CASE (condition-based — like if/else if), CASE in SELECT, WHERE, ORDER BY, CASE inside Aggregates (conditional counting), 4 practice problems (Easy → Crazy), 3 Try It Yourself challenges
Part 2: COALESCE, NULLIF & CAST	55 mins	COALESCE — first non-NULL value + practice, NULLIF — conditionally create NULL + practice, CAST & :: — type conversion + practice, Combined hard & crazy problems (all 4 tools), 2 Try It Yourself challenges
Quiz + Summary & Wrap-up	30 mins	Quick quiz (7 questions), Common mistakes gallery, Interview bank — conditional logic, Syntax Reference Card, Homework + Aggregates & Grouping preview

YOUR SQL JOURNEY

- ✓ SQL Intro
- ✓ Setup & RetailMart
- ✓ DDL & Data Types
- ✓ Constraints & DML
- ✓ Filtering & Sorting
- ✓ Scalar Functions (String & Date)
- 7** **Conditional Logic & Derived Columns**

← HERE

Memory Check

earlier sessions: SQL categories, PostgreSQL setup, RetailMart V3.

earlier sessions: DDL, Data Types, Constraints, DML.

SELECT, WHERE, AND/OR, LIKE, BETWEEN, IN, IS NULL, ORDER BY, LIMIT.

String functions (UPPER, TRIM, CONCAT, SPLIT_PART) + Date functions (NOW, EXTRACT, AGE, TO_CHAR, DATE_TRUNC).

Today: We can FIND and TRANSFORM data. Now we add DECISION-MAKING — different output based on conditions.

The Report Card Generator

A school needs to generate 2,000 report cards. Each student has raw marks (0-100). The system needs to:

- Convert marks to grades: A+ (90+), A (80-89), B (70-79), C (60-69), F (below 60)
- Handle absent students (NULL marks) — show 'ABSENT' not blank
- Display percentage as text with % symbol
- Avoid crash when total_students = 0 (division by zero)

This is EXACTLY what CASE WHEN, COALESCE, NULLIF, and CAST do.

From Story to SQL!

Convert marks to grades → `CASE WHEN marks >= 90 THEN 'A+'...`

Handle absent students → `COALESCE(marks, 'ABSENT')`

Format percentage → `CAST(marks AS TEXT) || '%'`

Avoid division by zero → `total / NULLIF(divisor, 0)`

`CASE WHEN` is the most powerful expression in SQL. It turns raw data into business-meaningful categories.

The Traffic Signal Analogy

Red → STOP. Yellow → SLOW DOWN. Green → GO.

A traffic signal checks the value and gives a result. CASE WHEN works the same way — it checks conditions and returns the first matching result. Like IF-ELSE in programming, but inside your SQL query.

DEFINITION

CASE WHEN

Technical:

CASE evaluates conditions top-to-bottom and returns the result of the FIRST matching condition. Two forms: Simple CASE (matches one expression against values) and Searched CASE (evaluates boolean conditions). If no match and no ELSE, returns NULL.

Layman's:

CASE WHEN is like a series of IF-ELSE rules. 'If the order is above 10,000 then VIP. If between 5,000-10,000 then Regular. Otherwise Budget.' SQL checks each rule in order and uses the first one that matches.

SYNTAX: Simple CASE

```
-- Simple CASE: match one column against values  
-- Like a switch statement in programming
```

```
SELECT column_name,  
       CASE column_name  
         WHEN 'value1' THEN 'result1'  
         WHEN 'value2' THEN 'result2'  
         WHEN 'value3' THEN 'result3'  
         ELSE 'default_result'  
       END AS alias_name  
FROM table_name;
```

```
-- Key rules:  
-- 1. Checks top-to-bottom, first match wins  
-- 2. ELSE catches everything not matched  
-- 3. Without ELSE, unmatched rows get NULL
```

PRO TIP

PRO TIP

Simple CASE is best when you are matching ONE column against EXACT values — like status codes, category names, or tier labels. If you need ranges (price > 1000) or multiple column checks, use Searched CASE instead.

SYNTAX: Searched CASE

```
-- Searched CASE: evaluate boolean conditions  
-- Like if/else if/else in programming
```

```
SELECT column_name,  
       CASE  
         WHEN condition1 THEN 'result1'  
         WHEN condition2 THEN 'result2'  
         WHEN condition3 THEN 'result3'  
         ELSE 'default_result'  
       END AS alias_name  
FROM table_name;
```

```
-- Key: conditions can use any operator  
-- >, <, >=, BETWEEN, LIKE, IN, IS NULL, etc.
```

Condition Order Matters!

The Mistake:

CASE WHEN price > 1000 THEN 'Budget' WHEN price > 10000 THEN 'Mid-Range' END — a product with price 50000 matches the FIRST condition (> 1000) and gets labeled 'Budget'!

The Fix:

Always put the MOST SPECIFIC condition first: WHEN price > 50000 ... WHEN price > 10000 ... WHEN price > 1000.
Check from highest to lowest.

CASE in WHERE — Conditional Filtering

```
-- Show only 'problem' orders based on rules
SELECT order_id, order_status, net_total,
       TO_CHAR(order_date, 'DD-Mon-YYYY') AS dt
FROM sales.orders
WHERE
  CASE
    WHEN order_status = 'Delivered'
      AND net_total > 20000
    THEN 'flag'
    WHEN order_status = 'Pending'
      AND order_date < '2025-06-01'
    THEN 'flag'
    WHEN order_status = 'Failed'
    THEN 'flag'
  END = 'flag'
ORDER BY order_date LIMIT 15;

-- CASE in WHERE = complex multi-condition
-- filtering that AND/OR alone cannot do
```

CASE in ORDER BY — Custom Sort Priority

```
-- Sort orders by urgency, not alphabetically
-- 'Pending' first, 'Delivered' last
SELECT order_id, order_status, net_total
FROM sales.orders
ORDER BY
    CASE order_status
        WHEN 'Pending' THEN 1
        WHEN 'Processing' THEN 2
        WHEN 'Out for Delivery' THEN 3
        WHEN 'Shipped' THEN 4
        WHEN 'Delivered' THEN 5
        WHEN 'Returned' THEN 6
        WHEN 'Cancelled' THEN 7
        WHEN 'Failed' THEN 8
        ELSE 9
    END,
    order_date ASC
LIMIT 20;
```

PRO TIP

PRO TIP

CASE in ORDER BY is extremely useful for dashboards. Users expect urgent items at the top — not alphabetical sorting. Assign numeric priorities (1 = most urgent) and ORDER BY that CASE expression ascending.

The Dashboard Single-Query Trick

Management wants ONE query that shows: total orders, delivered count, pending count, cancelled count. Without CASE inside aggregates, you would need 4 separate queries. With CASE inside COUNT, one query does it all.

SYNTAX: CASE Inside COUNT / SUM / AVG

-- Count only rows matching a condition

```
COUNT(CASE WHEN condition THEN 1 END)
```

-- Sum only matching amounts

```
SUM(CASE WHEN condition THEN amount END)
```

-- Average only matching values

```
AVG(CASE WHEN condition THEN value END)
```

-- How it works:

-- CASE returns 1 (or the value) for matches

-- CASE returns NULL for non-matches

-- COUNT/SUM/AVG skip NULLs automatically!

PRO TIP

PRO TIP

CASE inside aggregates is a Aggregates & Grouping preview — it combines conditional logic (Conditional Logic) with aggregate functions. This pattern appears in almost every analytics dashboard. Master it and you can build any summary report in a single query.

Easy: Customer Tier Labels

EASY

SCENARIO: The marketing team needs customer tiers displayed with friendly labels for the email campaign system.

TASK: Use Simple CASE on `customers.customers.tier` to map each tier to a display label: 'Gold' → 'Premium Member', 'Silver' → 'Valued Customer', 'Bronze' → 'Active Shopper', anything else → 'New Customer'. Show `customer_id`, name (INITCAP), tier, and the display label. Show 15 customers.

HINT: Two CASE forms exist — which one fits exact-value matching against a single column? Trap: don't forget an ELSE branch for unhandled values, or those rows will return NULL instead of a label.

Answer

✓ ANSWER

```
SELECT customer_id,  
       INITCAP(first_name) || ' ' ||  
       INITCAP(last_name) AS name,  
       tier,  
       CASE tier  
         WHEN 'Gold' THEN 'Premium Member'  
         WHEN 'Silver' THEN 'Valued Customer'  
         WHEN 'Bronze' THEN 'Active Shopper'  
         ELSE 'New Customer'  
       END AS display_tier  
FROM customers.customers  
LIMIT 15;
```

Medium: Employee Salary Band Classification

MEDIUM

SCENARIO: HR needs employees classified into salary bands for a compensation review: 'Executive' (salary > 100000), 'Senior' (60000-100000), 'Mid-Level' (30000-60000), 'Entry' (below 30000).

TASK: Use Searched CASE on stores.employees.salary. Show employee_id, name, role, salary, and salary_band. Order by salary descending. Show 20 employees.

HINT: Two CASE forms exist — which one supports range checks (>=, <)? Trap: order matters in WHEN clauses. The first match wins — if you put the lowest range first, higher values still fall into it.

Answer

 ANSWER

```
SELECT employee_id,  
       INITCAP(first_name) || ' ' ||  
       INITCAP(last_name) AS name,  
       role, salary,  
       CASE  
         WHEN salary > 100000 THEN 'Executive'  
         WHEN salary > 60000 THEN 'Senior'  
         WHEN salary > 30000 THEN 'Mid-Level'  
         ELSE 'Entry'  
       END AS salary_band  
FROM stores.employees  
ORDER BY salary DESC  
LIMIT 20;
```

Hard: Shipment Delivery Speed Rating

HARD

SCENARIO: The logistics team wants to rate deliveries by speed. If delivered within 3 days of shipping → 'Fast'. Within 7 days → 'Normal'. More than 7 days → 'Slow'. Not yet delivered → 'In Transit'.

TASK: Use Searched CASE on the day difference between delivered_date and shipped_date. Show shipment_id, courier_name, shipped_date, delivered_date, days taken, and the speed rating. Show 20 shipments.

HINT: delivered_date and shipped_date are DATE columns, so delivered_date - shipped_date is already an integer number of days — compare it directly. Trap: when delivered_date IS NULL, the subtraction returns NULL and your bucket logic fails — handle the NULL case as the FIRST WHEN in your CASE.

Answer

✓ ANSWER

```
SELECT shipment_id,  
       UPPER(courier_name) AS courier,  
       TO_CHAR(shipped_date, 'DD-Mon-YYYY')  
         AS shipped,  
       TO_CHAR(delivered_date, 'DD-Mon-YYYY')  
         AS delivered,  
       CASE  
         WHEN delivered_date IS NULL  
           THEN 'In Transit'  
         WHEN delivered_date - shipped_date <= 3  
           THEN 'Fast'  
         WHEN delivered_date - shipped_date <= 7  
           THEN 'Normal'  
         ELSE 'Slow'  
       END AS speed_rating  
FROM sales.shipments  
ORDER BY shipped_date DESC  
LIMIT 20;
```

Crazy: Payroll Tax Bracket Assignment

CRAZY

SCENARIO: The CFO wants to verify the income tax slabs on payroll.pay_slips. For each employee's gross salary, classify into Indian income-tax brackets: 0% ($\leq 2.5L$), 5% ($\leq 5L$), 20% ($\leq 10L$), 30% ($> 10L$). This is exactly the range-based CASE WHEN pattern.

TASK: From payroll.pay_slips, classify gross_salary into bracket labels. Show employee_id, salary_month, gross_salary, the expected tax_rate, and a column comparing 'Expected vs Actual' income_tax. Filter to most recent pay slips.

HINT: Range buckets are a Searched CASE — each WHEN tests an inequality, ORDERED from smallest to largest. Trap: Indian tax brackets use ANNUAL income but pay_slips show MONTHLY gross — multiply gross_salary by 12 before comparing to the annual thresholds, or your buckets misclassify everyone.

Answer

✓ ANSWER

```
SELECT
    employee_id,
    salary_month,
    gross_salary,
    gross_salary * 12 AS annual_estimate,
    CASE
        WHEN gross_salary * 12 <= 250000
            THEN '0% (Nil Bracket)'
        WHEN gross_salary * 12 <= 500000
            THEN '5% (Lower)'
        WHEN gross_salary * 12 <= 1000000
            THEN '20% (Middle)'
        ELSE '30% (Top)'
    END AS expected_bracket,
    income_tax AS actual_tax
FROM payroll.pay_slips
ORDER BY payment_date DESC
LIMIT 15;
```

Quick Fire Round!

Run each query on your RetailMart database. Modify the conditions. Add your own categories. Understanding comes from experimentation.

Challenge 1: Weekend vs Weekday Orders

```
-- Label each order as 'Weekend' or 'Weekday'
SELECT order_id, order_date,
       TO_CHAR(order_date, 'Day') AS day_name,
       CASE
         WHEN EXTRACT(DOW FROM order_date)
              IN (0, 6) THEN 'Weekend'
         ELSE 'Weekday'
       END AS day_type
FROM sales.orders
LIMIT 15;

-- YOUR TURN: Try adding 'Friday Evening'
-- as a separate category
```

Challenge 2: Product Availability Flag

```
-- Flag products based on inventory
SELECT p.product_id, p.product_name,
       p.price,
       CASE
         WHEN p.price > 50000 THEN 'Premium'
         WHEN p.price > 5000 THEN 'Standard'
         ELSE 'Economy'
       END AS segment
FROM products.products p
ORDER BY p.price DESC LIMIT 15;

-- YOUR TURN: Add a 'Luxury' tier
-- for products above 100000
```

Challenge 3: Employee Tenure Classification

```
-- Classify employees by tenure
SELECT employee_id,
       INITCAP(first_name) AS name, role,
       EXTRACT(YEAR FROM AGE(joining_date))
       AS years,
       CASE
         WHEN EXTRACT(YEAR FROM
           AGE(joining_date)) >= 5
         THEN 'Veteran'
         WHEN EXTRACT(YEAR FROM
           AGE(joining_date)) >= 2
         THEN 'Experienced'
         ELSE 'New Hire'
       END AS tenure_level
FROM stores.employees
ORDER BY joining_date LIMIT 15;
```

KEY TAKEAWAYS

Simple CASE: match one column against exact values (like switch). Best for status codes, category names.

Searched CASE: evaluate boolean conditions (like if/else). Best for ranges, complex logic.

Evaluates top-to-bottom — first match wins. Always include ELSE to avoid unexpected NULLs.

Can be used in SELECT, WHERE, ORDER BY — and inside aggregate functions (COUNT, SUM, AVG).

CASE inside aggregates = conditional counting/summing. Builds entire dashboards in one query.

CASE WHEN vs Application Logic

Q: 'Should business logic be in SQL or in the application?'

A: Simple categorizations (price tiers, status labels, salary bands) belong in SQL — they run faster and are reusable across reports. Complex multi-step logic (workflows, state machines, approval chains) belongs in the application. Rule of thumb: if a data analyst needs it in a report, put it in SQL.

The E-Commerce Dashboard Problem

Flipkart's dashboard shows customer profiles. Some have phone numbers, some don't. Some have middle names, some don't. Showing NULL on the screen looks broken — users think the app is glitchy.

COALESCE replaces NULLs with friendly defaults: 'Not Provided', 'N/A', or 0.

NULLIF does the reverse — creates NULL when you WANT to ignore a value (like dividing by zero).

CAST converts between types — number to text, text to date, etc.

COALESCE

Technical:

COALESCE(value1, value2, ..., valueN) returns the FIRST non-NULL value from the list. If all values are NULL, returns NULL. It is the SQL standard way to handle NULLs with fallback values.

Layman's:

COALESCE is like asking: 'Do you have a mobile number? No? Okay, home number? No? Fine, I will write N/A.' It tries each option in order and uses the first one that exists.

SYNTAX: COALESCE

```
-- Replace NULL with a default value
SELECT COALESCE(NULL, 'fallback');
    -- 'fallback'
SELECT COALESCE('hello', 'fallback');
    -- 'hello' (already non-NULL)

-- Multiple fallbacks – tries left to right
SELECT COALESCE(NULL, NULL, 'third');
    -- 'third'

-- With numbers
SELECT COALESCE(NULL, 0);    -- 0
SELECT COALESCE(42, 0);     -- 42
```

PRO TIP

COALESCE is better than CASE WHEN ... IS NULL for simple NULL defaults. COALESCE(phone, 'N/A') is cleaner than CASE WHEN phone IS NULL THEN 'N/A' ELSE phone END. Use COALESCE for simple NULL replacement, CASE WHEN for complex conditions.

Easy: Clean Customer Display

EASY

SCENARIO: The customer profile page shows ugly 'null' text for missing phone numbers. The frontend team wants 'Not Provided' instead.

TASK: Show customer_id, name (INITCAP), phone with COALESCE fallback to 'Not Provided', email, and tier with COALESCE fallback to 'Standard'. Show 15 rows.

HINT: Which function returns the first non-NULL value from a list of arguments? Trap: the fallback value must be the SAME data type as the column — wrapping a numeric column with COALESCE(col, 'N/A') causes a type-cast error.

Answer

✓ ANSWER

```
SELECT customer_id,  
       INITCAP(first_name) || ' ' ||  
       INITCAP(last_name) AS name,  
       COALESCE(phone, 'Not Provided')  
       AS phone,  
       email,  
       COALESCE(tier, 'Standard') AS tier  
FROM customers.customers  
LIMIT 15;
```

Medium: Multi-Fallback Contact Info

MEDIUM

SCENARIO: The notification system needs one primary contact method per customer. Try phone first, then email, then show 'No Contact Info'. Also classify tier with a CASE label.

TASK: Show customer_id, name, phone, email, a primary_contact column using COALESCE with multiple fallbacks, and a tier_label using CASE. Show 15 customers WHERE phone IS NULL.

HINT: COALESCE accepts MORE than 2 arguments — it walks the list left-to-right until it finds a non-NULL. Trap: order matters in business logic — the highest-priority source must come first, since later arguments only get checked if everything before them is NULL.

Answer

✓ ANSWER

```
SELECT customer_id,  
       INITCAP(first_name) || ' ' ||  
       INITCAP(last_name) AS name,  
       phone, email,  
       COALESCE(phone, email,  
        'No Contact Info')  
       AS primary_contact,  
       CASE tier  
         WHEN 'Gold' THEN 'Premium'  
         WHEN 'Silver' THEN 'Valued'  
         WHEN 'Bronze' THEN 'Active'  
         ELSE 'Standard'  
       END AS tier_label  
FROM customers.customers  
WHERE phone IS NULL  
LIMIT 15;
```

NULLIF

Technical:

NULLIF(value1, value2) returns NULL if value1 equals value2, otherwise returns value1. Its primary use is preventing division-by-zero errors.

Layman's:

NULLIF says: 'If this value is X, pretend it does not exist (make it NULL).' Most used for: NULLIF(count, 0) before dividing — because dividing by NULL gives NULL (safe), but dividing by zero crashes your query.

SYNTAX: NULLIF

```
-- Returns NULL if both values are equal
SELECT NULLIF(10, 10); -- NULL (equal!)
SELECT NULLIF(10, 5);  -- 10 (not equal)
SELECT NULLIF('abc', 'abc'); -- NULL
SELECT NULLIF('abc', 'xyz'); -- 'abc'

-- The killer use case: safe division
-- BAD:  SELECT 100 / 0; -- ERROR! Crash!
-- GOOD: SELECT 100 / NULLIF(0, 0); -- NULL

-- In practice:
SELECT total_revenue / NULLIF(total_orders, 0)
   AS avg_order_value
FROM some_summary;
-- If total_orders = 0, returns NULL not error
```

PRO TIP

PRO TIP

NULLIF + COALESCE combo: `COALESCE(revenue / NULLIF(orders, 0), 0)`. This says: divide revenue by orders (safely — if orders is 0 return NULL), then if the result is NULL, show 0 instead. Two safety nets in one expression.

Easy: Safe Product Profit Margin

EASY

SCENARIO: The pricing team wants profit margin for each product: $(\text{price} - \text{cost_price}) / \text{price} * 100$. But some products have price = 0 (free samples), causing division-by-zero crashes.

TASK: Show product_id, product_name, price, cost_price, and profit_margin_pct using NULLIF(price, 0) to prevent crashes. Round to 1 decimal. Show 15 products ordered by price DESC.

HINT: Division by zero crashes the query — which function turns a specific value (like 0) into NULL? Trap: NULL propagates through math, so the result becomes NULL — that's safer than crashing, but you may still want to wrap the whole expression in COALESCE to return 0 instead.

Answer

✓ ANSWER

```
SELECT product_id, product_name,  
       price, cost_price,  
       ROUND((price - cost_price) /  
             NULLIF(price, 0) * 100, 1)  
       AS profit_margin_pct  
FROM products.products  
ORDER BY price DESC  
LIMIT 15;
```

Medium: Margin with Safe Default

MEDIUM

SCENARIO: Extend the margin report — free products (price=0) should show 0% instead of NULL. Also add a CASE to label margins: 'High' (>50%), 'Normal' (20-50%), 'Low' (<20%).

TASK: Show product_id, product_name, price, margin_pct (NULLIF + COALESCE for default 0), and margin_label (CASE on the margin). Show 20 products.

HINT: You need TWO defenses — one to handle a zero price, another to handle the resulting NULL. Trap: you can't reference a SELECT alias inside another SELECT expression on the same level — you'll have to repeat the calculation in the CASE, or wrap the whole thing in a derived table.

Answer



ANSWER

```
SELECT product_id, product_name,  
       price,  
       COALESCE(ROUND(  
         (price - cost_price) /  
         NULLIF(price, 0) * 100, 1  
       ), 0) AS margin_pct,  
       CASE  
         WHEN COALESCE(ROUND(  
           (price - cost_price) /  
           NULLIF(price, 0) * 100, 1  
         ), 0) > 50 THEN 'High Margin'  
         WHEN COALESCE(ROUND(  
           (price - cost_price) /  
           NULLIF(price, 0) * 100, 1  
         ), 0) > 20 THEN 'Normal Margin'  
         ELSE 'Low Margin'  
       END AS margin_label  
FROM products.products  
ORDER BY price DESC  
LIMIT 20;
```


CAST

Technical:

CAST(expression AS target_type) converts a value from one data type to another. PostgreSQL also supports the shorthand :: operator. Common conversions: TEXT ↔ INTEGER, TEXT ↔ DATE, NUMERIC ↔ INTEGER, TIMESTAMP ↔ DATE.

Layman's:

CAST is like changing the format of data. Convert the number 42 to the text '42' so you can concatenate it. Convert the text '2025-01-15' to a proper DATE so you can do date math. It does not change the stored data — only the output.

SYNTAX: CAST and :: (PostgreSQL shorthand)

```
-- SQL Standard: CAST(value AS type)
SELECT CAST(42 AS TEXT);           -- '42'
SELECT CAST('42' AS INTEGER);     -- 42
SELECT CAST('2025-01-15' AS DATE);
SELECT CAST(99.99 AS INTEGER);    -- 99 (truncates!)

-- PostgreSQL shorthand: value::type
SELECT 42::TEXT;                  -- '42'
SELECT '42'::INTEGER;            -- 42
SELECT '2025-01-15'::DATE;
SELECT 99.99::INTEGER;           -- 99

-- Both do the same thing
-- :: is shorter and preferred in PostgreSQL
```

CAST Truncation — Silent Data Loss

The Mistake:

`CAST(99.99 AS INTEGER)` returns 99 — the decimal part is silently dropped. No error, no warning. Your report shows wrong numbers.

The Fix:

Use `ROUND(value)` before casting: `ROUND(99.99)::INTEGER` gives 100. Or use `NUMERIC` for precision: `value::NUMERIC(10,2)`. Always test `CAST` with sample data before applying to reports.

Easy: Product Codes with Type Conversion

EASY

SCENARIO: The frontend team needs product IDs displayed as 'PROD-1234' and formatted prices as 'Rs. X,XXX' for the catalog page. Currently `product_id` is an integer — concatenating it directly with text causes a type error.

TASK: Show `product_id`, 'PROD-' || `product_id::TEXT` as `product_code`, `product_name`, and 'Rs. ' || `TO_CHAR(price, 'FM99,99,999')` as `display_price`. Show 15 products.

HINT: Two conversions needed — one to allow concatenation of a number with text, one to format the price with thousand separators. Trap: || on a number+string without an explicit cast throws a type error; integers need to become TEXT first.

Answer

✓ ANSWER

```
SELECT product_id,  
       'PROD-' || product_id::TEXT  
       AS product_code,  
       product_name,  
       'Rs. ' || TO_CHAR(price,  
       'FM99,99,999') AS display_price  
FROM products.products  
ORDER BY price DESC  
LIMIT 15;
```

Medium: Formatted Order Summary Line

MEDIUM

SCENARIO: Build a formatted order report. Each line should look like: 'Order #12345 — Rs. 15,000 — 15-Jan-2025 — Delivered'. This requires converting numbers and dates to text for concatenation.

TASK: Build a single summary_line column that concatenates order_id (as text), formatted amount (TO_CHAR), formatted date (TO_CHAR), and status (INITCAP). Show 15 orders from 2025.

HINT: Three formatting decisions — integer to text, number to currency string, date to a human-readable format. Trap: TO_CHAR with numbers adds a leading space for the sign by default; a 'FM' prefix in the format string suppresses that padding.

Answer

 ANSWER

```
SELECT
    'Order #' || order_id::TEXT ||
    ' - Rs. ' || TO_CHAR(net_total,
        'FM99,99,999') ||
    ' - ' || TO_CHAR(order_date,
        'DD-Mon-YYYY') ||
    ' - ' || INITCAP(order_status)
    AS summary_line
FROM sales.orders
WHERE order_date >= '2025-01-01'
ORDER BY order_date DESC
LIMIT 15;
```

Hard: Complete Customer Health Report

HARD

SCENARIO: Build a customer profile card combining ALL Conditional Logic tools. For each customer: a customer code, clean name, phone status, account age category, and tier label.

TASK: Use CAST for customer code ('CUST-' || id::TEXT), COALESCE for phone ('N/A' if NULL), CASE for age category (5+ yrs = 'Veteran', 2-4 = 'Loyal', else = 'New'), and CASE for tier label. Show 20 customers ordered by registration_date.

HINT: AGE() produces an interval between dates — pair it with EXTRACT to pull out just the year part. Trap: AGE() with one argument compares to TODAY; with two arguments it compares between them. Pick the form that matches your business question.

Answer (1/2)

✓ ANSWER

```
SELECT
  'CUST-' || customer_id::TEXT
    AS cust_code,
  INITCAP(first_name) || ' ' ||
    INITCAP(last_name) AS name,
  COALESCE(phone, 'N/A') AS phone,
  CASE
    WHEN EXTRACT(YEAR FROM
      AGE(registration_date)) >= 5
      THEN 'Veteran'
    WHEN EXTRACT(YEAR FROM
      AGE(registration_date)) >= 2
      THEN 'Loyal'
    ELSE 'New'
  END AS age_category,
  CASE tier
    WHEN 'Gold' THEN 'Premium'
    WHEN 'Silver' THEN 'Valued'
    WHEN 'Bronze' THEN 'Active'
    ELSE 'Standard'
  END AS tier_label
FROM customers.customers
```

Answer (2/2)

✓ ANSWER

```
ORDER BY registration_date  
LIMIT 20;
```

Crazy: Loyalty Tier Mismatch Detection

CRAZY

SCENARIO: The CMO suspects some loyalty members are in the WRONG tier. Each tier has a min_points threshold. A member's current points_balance should map to the correct tier — but maybe doesn't. Use CASE on points_balance to derive the SHOULD-BE tier; compare to actual.

TASK: Join loyalty.members to loyalty.tiers and customers. For each member, show name, points balance, current_tier (actual), should_be_tier (derived via CASE on points), and a 'mismatch' flag. Use NULLIF to dodge division by zero in a 'closeness to next tier' percentage.

HINT: This combines all four Conditional Logic tools — CASE for derived tier, COALESCE for missing tier names, NULLIF for safe division, CAST for clean string output. Trap: comparing two strings ('Gold' vs 'Gold') seems easy but case sensitivity bites — wrap both sides with UPPER or LOWER when matching tier names.

Answer (1/2)

✓ ANSWER

```
SELECT
  c.customer_id,
  c.first_name || ' ' || c.last_name
    AS member_name,
  m.points_balance,
  COALESCE(t.tier_name,
    'Unassigned') AS current_tier,
  CASE
    WHEN m.points_balance >= 5000
      THEN 'Platinum'
    WHEN m.points_balance >= 2000
      THEN 'Gold'
    WHEN m.points_balance >= 500
      THEN 'Silver'
    ELSE 'Bronze'
  END AS should_be_tier,
  CASE
    WHEN UPPER(COALESCE(t.tier_name, ''))
      = UPPER(CASE
        WHEN m.points_balance >= 5000
          THEN 'Platinum'
        WHEN m.points_balance >= 2000
```

Answer (2/2)

✓ ANSWER

```
        THEN 'Gold'
      WHEN m.points_balance >= 500
        THEN 'Silver'
      ELSE 'Bronze'
    END) THEN 'OK'
  ELSE 'MISMATCH'
END AS status
FROM loyalty.members m
JOIN customers.customers c
  ON m.customer_id = c.customer_id
LEFT JOIN loyalty.tiers t
  ON m.tier_id = t.tier_id
ORDER BY m.points_balance DESC LIMIT 20;
```

Hands-On Practice!

Run each query on your RetailMart database. Modify the values, add your own columns. The best way to learn is to experiment.

Challenge 1: Employee Directory with Defaults

```
-- Build a clean employee display
SELECT employee_id,
       INITCAP(first_name) || ' ' ||
       INITCAP(last_name) AS name,
       COALESCE(role, 'Unassigned') AS role,
       'EMP-' || LPAD(employee_id::TEXT,
       5, '0') AS emp_code,
       salary::TEXT || '/month' AS salary_display
FROM stores.employees LIMIT 15;

-- YOUR TURN: Add a CASE for salary band
```

Challenge 2: Product Price Display

```
-- Format product prices for catalog
SELECT product_id,
       product_name,
       'Rs. ' || TO_CHAR(price,
       'FM99,99,999') AS display_price,
       CASE
         WHEN cost_price > 0 THEN
           ROUND((price - cost_price) /
           NULLIF(price, 0) * 100)::TEXT
           || '% margin'
         ELSE 'N/A'
       END AS margin_info
FROM products.products
WHERE price > 5000
ORDER BY price DESC LIMIT 15;
```


Answer Before Looking!

For each question, try to answer BEFORE scrolling to the answer. This builds interview confidence.

Quiz 1: What does this return?

EASY

ANSWER: 'B' — CASE checks top-to-bottom. $15 > 20$ is FALSE. $15 > 10$ is TRUE → returns 'B'. It STOPS checking — never reaches 'C' even though $15 > 5$ is also TRUE.

Question



ANSWER

```
SELECT CASE
  WHEN 15 > 20 THEN 'A'
  WHEN 15 > 10 THEN 'B'
  WHEN 15 > 5  THEN 'C'
END;
```

Quiz 2: What does this return?

EASY

ANSWER: 'hello' — COALESCE returns the FIRST non-NULL value. Skips the two NULLs, finds 'hello', and stops. Never looks at 'world'.

Question



ANSWER

```
SELECT COALESCE(NULL, NULL, 'hello', 'world');
```

Quiz 3: What does this return?

MEDIUM

ANSWER: NULL — `NULLIF(0, 0)` returns NULL because both values are equal. $100 / \text{NULL} = \text{NULL}$. Without `NULLIF`, $100 / 0$ would crash with a division-by-zero error.

Question



ANSWER

```
SELECT 100 / NULLIF(0, 0);
```

Quiz 4: What does this return?

MEDIUM

ANSWER: 'Not Equal' — `NULL = NULL` is UNKNOWN (not TRUE). In SQL, nothing equals NULL, not even NULL itself. To check for NULL, use `IS NULL`, not `= NULL`.

Question



ANSWER

```
SELECT CASE WHEN NULL = NULL THEN 'Equal'  
          ELSE 'Not Equal' END;
```

Quiz 5: Spot the Bug

HARD

ANSWER: Two bugs! (1) Condition order is wrong — a 60000 product matches the FIRST condition (> 1000) and gets 'Budget'. Fix: put largest value first. (2) No ELSE — products with price ≤ 1000 get NULL. Fix: add ELSE 'Value'.

Question

 ANSWER

```
SELECT product_name, price,  
       CASE  
         WHEN price > 1000 THEN 'Budget'  
         WHEN price > 10000 THEN 'Mid'  
         WHEN price > 50000 THEN 'Premium'  
       END AS tier  
FROM products.products;
```

Quiz 6: What does this return?

HARD

ANSWER: It returns the price as TEXT (e.g. '4999'). Since price IS NOT NULL (WHERE clause), COALESCE always returns the first argument. The ::TEXT cast is needed because 'N/A' is TEXT — all COALESCE args must match types.

Question



ANSWER

```
SELECT COALESCE(price::TEXT, 'N/A')  
FROM products.products  
WHERE price IS NOT NULL LIMIT 1;
```

Quiz 7: Predict the Output

CRAZY

ANSWER: 42. First arg is NULL → skip. Second arg: NULLIF(5, 5) returns NULL because both values are equal → skip. Third arg: 42 is not NULL → return 42. This tests understanding of both COALESCE and NULLIF together.

Question

 ANSWER

```
SELECT  
  COALESCE(NULL, NULLIF(5, 5), 42);
```

Forgetting ELSE — Silent NULLs

The Mistake:

CASE WHEN price > 10000 THEN 'Expensive' END — products with price <= 10000 get NULL instead of a label. No error, just blank cells in your report.

The Fix:

Always add ELSE: CASE WHEN price > 10000 THEN 'Expensive' ELSE 'Affordable' END. Without ELSE, unmatched rows silently become NULL.

NULL = NULL is NOT TRUE

The Mistake:

CASE WHEN phone = NULL THEN 'Missing' END — this NEVER matches because NULL = NULL is UNKNOWN, not TRUE.
No rows ever get labeled 'Missing'.

The Fix:

Use IS NULL: CASE WHEN phone IS NULL THEN 'Missing' ELSE phone END. Or better yet: COALESCE(phone, 'Missing').

COALESCE Type Mismatch

The Mistake:

`COALESCE(price, 'N/A')` — `price` is `NUMERIC`, `'N/A'` is `TEXT`. PostgreSQL throws a type error because all `COALESCE` arguments must be the same type.

The Fix:

Cast to the same type: `COALESCE(price::TEXT, 'N/A')`. Or use `CASE WHEN`: `CASE WHEN price IS NULL THEN 'N/A' ELSE price::TEXT END`.

CASE WHEN with OR vs Multiple WHENs

The Mistake:

Using OR incorrectly: `CASE WHEN status = 'Delivered' OR 'Shipped' THEN 'Done' END` — this is a syntax error. 'Shipped' alone is not a boolean expression.

The Fix:

Either: `CASE WHEN status = 'Delivered' OR status = 'Shipped' THEN 'Done' END`. Or cleaner: `CASE WHEN status IN ('Delivered', 'Shipped') THEN 'Done' END`.

Simple CASE vs Searched CASE

Q: 'What is the difference between Simple CASE and Searched CASE?'

A: Simple CASE matches ONE expression against values: CASE status WHEN 'Active' THEN ... Good for equality checks on a single column. Searched CASE evaluates boolean conditions: CASE WHEN price > 1000 THEN ... Good for ranges, comparisons across columns, and complex logic.

COALESCE vs IFNULL vs NVL

Q: 'What is COALESCE and how does it differ from IFNULL?'

A: COALESCE is SQL standard — works in PostgreSQL, MySQL, SQL Server, Oracle. It accepts N arguments and returns the first non-NULL. IFNULL is MySQL-only (2 arguments). NVL is Oracle-only. ISNULL is SQL Server-only. Always use COALESCE for portability and flexibility.

Prevent Division by Zero

Q: 'How do you handle division by zero in SQL?'

A: Use NULLIF: `revenue / NULLIF(order_count, 0)`. NULLIF returns NULL when `order_count` is 0, and anything divided by NULL gives NULL instead of an error. Optionally wrap with COALESCE for a default: `COALESCE(revenue / NULLIF(order_count, 0), 0)`.

CASE WHEN NULL Gotcha

Q: 'What does CASE WHEN column = NULL return?'

A: It NEVER matches because `NULL = NULL` evaluates to `UNKNOWN` (not `TRUE`). To check for `NULL` inside `CASE`, use `IS NULL`: `CASE WHEN column IS NULL THEN 'missing' END`. This is one of the most common SQL interview traps.

Conditional Aggregation

Q: 'How do you count only certain rows within an aggregate?'

A: Use CASE inside COUNT: COUNT(CASE WHEN order_status = 'Delivered' THEN 1 END). For sums: SUM(CASE WHEN category = 'Electronics' THEN amount END). Non-matching rows return NULL, which is skipped by aggregate functions. This creates pivot-style reports in a single query.

CAST vs TO_CHAR vs ::TEXT

Q: 'When do you use CAST vs TO_CHAR for converting numbers to text?'

A: CAST(42 AS TEXT) or 42::TEXT gives '42' — plain conversion. TO_CHAR(42000, 'FM99,999') gives '42,000' — formatted with commas. Use CAST/:: for simple type conversion. Use TO_CHAR when you need specific formatting (commas, decimals, currency).

CASE WHEN vs DECODE

Q: 'What is DECODE and how does it compare to CASE?'

A: DECODE is Oracle-only. It works like Simple CASE: `DECODE(status, 'A', 'Active', 'I', 'Inactive', 'Unknown')`. CASE WHEN is SQL standard — works in PostgreSQL, MySQL, SQL Server, Oracle. Always prefer CASE WHEN for portability. Plus, CASE supports conditions (ranges, LIKE, IN) that DECODE cannot do.

Real Interview: Classify and Count

Q: 'Write a query to show total orders, and what percentage are high-value (> 10000), medium (5000-10000), and low (< 5000).'

A: This combines Conditional Logic + Aggregates & Grouping: `ROUND(100.0 * COUNT(CASE WHEN net_total > 10000 THEN 1 END) / COUNT(*), 1) AS high_pct`. Repeat for medium and low tiers. The interviewer tests if you know CASE inside aggregates — the most powerful pattern in analytics SQL.

KEY TAKEAWAYS

CASE WHEN:

1. Simple CASE: match one column against fixed values (like switch). Good for status mapping.
2. Searched CASE: evaluate boolean conditions (like if/else). Good for ranges, complex logic.
3. Evaluates top-to-bottom — first match wins. Always include ELSE.
4. Can be used in SELECT, WHERE, ORDER BY, and inside aggregates (conditional counting/summing).

NULL HANDLING & TYPE CONVERSION:

5. COALESCE(a, b, c) — returns first non-NULL. Essential for clean reports.
6. NULLIF(a, b) — returns NULL if a = b. Essential for safe division.
7. CAST(x AS TYPE) or x::TYPE — convert between data types.
8. Pattern: SUM(CASE WHEN condition THEN amount END) = conditional aggregation.

Conditional Logic & Derived Columns

-- Simple CASE

```
CASE column
  WHEN 'val1' THEN 'result1'
  WHEN 'val2' THEN 'result2'
  ELSE 'default'
END AS alias
```

-- Searched CASE

```
CASE
  WHEN condition1 THEN 'result1'
  WHEN condition2 THEN 'result2'
  ELSE 'default'
END AS alias
```

-- CASE in Aggregates

```
COUNT(CASE WHEN status = 'Done' THEN 1 END)
SUM(CASE WHEN cat = 'X' THEN amt END)
AVG(CASE WHEN price > 1000 THEN price END)
```

Conditional Logic & Derived Columns

-- COALESCE

```
COALESCE(phone, 'N/A')    -- NULL → 'N/A'  
COALESCE(a, b, c, 0)      -- first non-NULL
```

-- NULLIF

```
NULLIF(count, 0)          -- 0 → NULL  
total / NULLIF(count, 0)  -- safe division
```

-- CAST & ::

```
CAST(price AS INTEGER)  
price::INTEGER            -- PG shorthand  
'2025-01-15'::DATE  
42::TEXT                  -- for concatenation
```

CHALLENGE

Conditional Logic Take-Home Challenge (1/3)

Build these queries on RetailMart V3:

Conditional Logic Take-Home Challenge (2/3)

- 1.: Classify all orders as 'High Value' (> 10000), 'Medium' (5000-10000), 'Low' (< 5000). Show order_id, net_total, and the classification.
- 2.: Map order statuses to user-friendly labels using Simple CASE. Sort by custom priority (Pending first) using CASE in ORDER BY.
- 3.: Show all customers with COALESCE on phone ('N/A') and tier ('Standard'). Include name (INITCAP) and account age.
- 4.: Build a customer profile: 'CUST-' || id, name, phone status (CASE: Has Phone / Missing), tier label (CASE), and registration year.
- 5.: Create a single-row dashboard: total orders, delivered count, pending count, cancelled count, avg delivered value — using CASE inside aggregates.
- 6.: Build an order summary line: 'Order #ID — Rs. X,XXX — DD-Mon-YYYY — Status' using CAST, TO_CHAR, and || concatenation.

CHALLENGE

Conditional Logic Take-Home Challenge (3/3)

Push your SQL file to GitHub!

<https://github.com/starlordali4444/postgresql>

WHAT YOU ACHIEVED TODAY

Your queries can now THINK. CASE WHEN adds decision-making. COALESCE handles NULLs gracefully. NULLIF prevents crashes. CAST converts between types. Combined with earlier sessions skills, you can now build real analytics dashboards.

Tomorrow: Aggregates — COUNT, SUM, AVG, MIN, MAX, GROUP BY, HAVING. You will learn to SUMMARIZE entire tables into business reports. Combined with today's CASE WHEN, you will build powerful pivot-style reports.

Coming Tomorrow

How many customers per city? What is the total revenue per month? Which store has the highest average order value?

GROUP BY + aggregate functions (COUNT, SUM, AVG, MIN, MAX) turn millions of rows into concise business reports.

HAVING filters groups (WHERE filters rows, HAVING filters groups).

Aggregates & Grouping is where SQL becomes a REPORTING POWERHOUSE. Combined with CASE WHEN, you will create dashboards that executives rely on.